

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on
        minibatches of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i]
        = c means
          that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # Initialize the loss to zero.
        loss = 0.0

        # =====
        #
        # YOUR CODE HERE:
        #   Calculate the normalized softmax loss. Store it as the
        variable loss.
        #   (That is, calculate the sum of the losses of all the training
        #   set margins, and then normalize the loss by the number of
        #   training examples.)
        # =====
        #
        num_train = X.shape[0]
        num_classes = self.W.shape[0]

        for i in np.arange(num_train):

```

```

        cum = 0
        aix = self.W[y[i]].dot(X[i])
        logk = -np.max(aix)

        for j in np.arange(num_classes):
            cum += np.exp(self.W[j].dot(X[i])) + logk

        loss += np.log(cum) - (aix + logk)

    loss /= num_train

    # =====
#
# END YOUR CODE HERE
# =====
#

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the
    gradient.

    Output: grad -- a matrix of the same dimensions as W
    containing
        the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # =====
#
# YOUR CODE HERE:
#     Calculate the softmax loss and the gradient. Store the
gradient
#     as the variable grad.
# =====
#

    num_train = X.shape[0]
    num_classes = self.W.shape[0]

    for i in np.arange(num_train):
        cum = 0
        aix = self.W[y[i]].dot(X[i])
        logk = -np.max(aix)

```

```

        for j in np.arange(num_classes):
            cum += np.exp(self.W[j].dot(X[i]) + logk)

        loss += np.log(cum) - (aix + logk)
        grad[y[i]] -= X[i]

    loss /= num_train
    grad /= num_train

    # =====
#
# END YOUR CODE HERE
# =====
#

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) /
        (abs(grad_numerical) + abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' %
              (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # =====
#

```

```

    # YOUR CODE HERE:
        # Calculate the softmax loss and gradient WITHOUT any for
loops.
    # =====
#

    num_train = X.shape[0]
    selector = np.arange(num_train), y

    scores = X.dot(self.W.T)
    scores -= np.max(scores)
    loss = np.sum(np.log(np.sum(np.exp(scores).T, axis=0)) -
scores[selector]) / num_train

    e_exp_a = np.exp(scores)
    sums = np.sum(e_exp_a, axis=1)
    temp = e_exp_a/sums[:, np.newaxis]
    temp[selector] -= 1

    grad = temp.T.dot(X) / num_train

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there
are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i]
= c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each
step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training
iteration.

```

```

"""
num_train, dim = X.shape
num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where
K is number of classes

self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) #
initializes the weights of self.W

# Run stochastic gradient descent to optimize W
loss_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    #
===== #
    # YOUR CODE HERE:
    # Sample batch_size elements from the training data for use in
    # gradient descent. After sampling,
    # - X_batch should have shape: (dim, batch_size)
    # - y_batch should have shape: (batch_size,)
    # The indices should be randomly generated to reduce
correlations
    # in the dataset. Use np.random.choice. It's okay to
sample with
    # replacement.
    #
===== #

    indices = np.random.choice(num_train, batch_size)
    X_batch, y_batch = X[indices], y[indices]

    #
===== #
    # END YOUR CODE HERE
    #
===== #

    # evaluate loss and gradient
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
    loss_history.append(loss)

    #
===== #
    # YOUR CODE HERE:
    # Update the parameters, self.W, with a gradient step
    #
===== #

```

```

        self.W -= learning_rate * grad

        #
===== #
        # END YOUR CODE HERE
        #
===== #

        if verbose and it % 100 == 0:
            print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

        return loss_history

    def predict(self, X):
        """
        Inputs:
        - X: N x D array of training data. Each row is a D-dimensional
point.

        Returns:
        - y_pred: Predicted labels for the data in X. y_pred is a 1-
dimensional
            array of length N, and each element is an integer giving the
predicted
            class.
        """
        y_pred = np.zeros(X.shape[1])
        # =====
#
        # YOUR CODE HERE:
        #   Predict the labels given the training data.
        # =====
#

        scores = X.dot(self.W.T)
        y_pred = np.argmax(scores, axis=1)

        # =====
#
        # END YOUR CODE HERE
        # =====
#

        return y_pred

```