

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University,
and modified for ECE C147/C247 at UCLA.
"""
class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has
        shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on
        minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i]
        = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0

        for i in np.arange(num_train):
            # =====
            # YOUR CODE HERE:
            # Calculate the normalized SVM loss, and store it as
            'loss'.
            # (That is, calculate the sum of the losses of all the training

```

```

        # set margins, and then normalize the loss by the number of
        # training examples.)
        # =====
#
    for j in np.arange(num_classes):
        if y[i] != j:
            loss += max(0, 1 + self.W[j].dot(X[i]) -
self.W[y[i]].dot(X[i]))

    loss /= num_train

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the
    gradient.

    Output: grad -- a matrix of the same dimensions as W
    containing
        the gradient of the loss with respect to W.
    """

    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros_like(self.W)

    for i in np.arange(num_train):
        # =====
#
        # YOUR CODE HERE:
        # Calculate the SVM loss and the gradient. Store the
        gradient in
        # the variable grad.
        # =====
#

        for j in np.arange(num_classes):
            if y[i] != j:
                zj = 1 + self.W[j].dot(X[i]) - self.W[y[i]].dot(X[i])
                loss += max(0, zj)
                grad[j] += 0 if zj <= 0 else X[i]
                grad[y[i]] += 0 if zj <= 0 else -X[i]

```

```

# =====
#
# END YOUR CODE HERE
# =====
#

    loss /= num_train
    grad /= num_train

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) /
        (abs(grad_numerical) + abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' %
        (grad_numerical, grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and ouptuts as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

# =====
#
# YOUR CODE HERE:
#     Calculate the SVM loss WITHOUT any for loops.
# =====
#

```

```

    scores = X.dot(self.W.T)
    num_train = X.shape[0]

    losses = np.ones_like(scores) + scores -
scores[np.arange(num_train), y].reshape(scores.shape[0], 1)
    losses[losses < 0] = 0
    losses[np.arange(num_train), y] = 0
    loss = np.sum(losses) / num_train

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    #
===== #
    # YOUR CODE HERE:
        # Calculate the SVM grad WITHOUT any for loops.
    # =====
#

    selector = losses.T
    selector[selector > 0] = 1
    selector[y, np.arange(num_train)] = np.sum(selector, axis=0) *
(-1)
    grad = selector.dot(X) / num_train

    # =====
#
    # END YOUR CODE HERE
    # =====
#

    return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
        batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there
are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i]
= c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing

```

- batch_size: (integer) number of training examples to use at each step.
- verbose: (boolean) If true, print progress during optimization.

Outputs:

A list containing the value of the loss function at each training iteration.

```

"""
num_train, dim = X.shape
num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where
K is number of classes

```

```

self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) #
initializes the weights of self.W

```

```

# Run stochastic gradient descent to optimize W
loss_history = []

```

```

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

```

```

#

```

```

===== #
# YOUR CODE HERE:
# Sample batch_size elements from the training data for use in
# gradient descent. After sampling,
# - X_batch should have shape: (dim, batch_size)
# - y_batch should have shape: (batch_size,)
# The indices should be randomly generated to reduce
correlations
# in the dataset. Use np.random.choice. It's okay to
sample with
# replacement.
#

```

```

===== #

```

```

indices = np.random.choice(num_train, batch_size)
X_batch, y_batch = X[indices], y[indices]

```

```

#

```

```

===== #
# END YOUR CODE HERE
#

```

```

===== #

```

```

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

```

```

#
===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step
#
===== #

self.W -= learning_rate * grad

#
===== #
# END YOUR CODE HERE
#
===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional
point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-
dimensional
    array of length N, and each element is an integer giving the
predicted
    class.
    """
    y_pred = np.zeros(X.shape[1])

# =====
#
# YOUR CODE HERE:
#   Predict the labels given the training data with the parameter
self.W.
# =====
#

scores = X.dot(self.W.T)
y_pred = np.argmax(scores, axis=1)

# =====
#
# END YOUR CODE HERE

```

```
# =====  
#  
return y_pred
```