

```

import numpy as np
import matplotlib.pyplot as plt

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu). It has been modified in various areas for use
in the
ECE 239AS class at UCLA. This includes the descriptions of what code
to
implement as well as some slight potential changes in variable names
to be
consistent with class nomenclature. We thank Justin Johnson & Serena
Yeung for
permission to use this code. To see the original version, please
visit
cs231n.stanford.edu.
"""

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network. The net has an input
    dimension of
    N, a hidden layer dimension of H, and performs classification over C
    classes.
    We train the network with a softmax loss function and L2
    regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the
    first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer -
    softmax

    The outputs of the second fully-connected layer are the scores for
    each class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
        Initialize the model. Weights are initialized to small random
        values and
        biases are initialized to zero. Weights and biases are stored in
        the
        variable self.params, which is a dictionary with the following
        keys:

        W1: First layer weights; has shape (H, D)
        b1: First layer biases; has shape (H,)

```

W2: Second layer weights; has shape (C, H)
b2: Second layer biases; has shape (C,)

Inputs:

- input_size: The dimension D of the input data.
- hidden_size: The number of neurons H in the hidden layer.
- output_size: The number of classes C.

"""

```
self.params = {}  
self.params['W1'] = std * np.random.randn(hidden_size, input_size)  
self.params['b1'] = np.zeros(hidden_size)  
self.params['W2'] = std * np.random.randn(output_size,  
hidden_size)  
self.params['b2'] = np.zeros(output_size)
```

```
def loss(self, X, y=None, reg=0.0):
```

"""

Compute the loss and gradients for a two layer fully connected
neural
network.

Inputs:

- X: Input data of shape (N, D). Each X[i] is a training sample.
- y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
an integer in the range $0 \leq y[i] < C$. This parameter is optional; if it
is not passed then we only return scores, and if it is passed then we
instead return the loss and gradients.
- reg: Regularization strength.

Returns:

If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
the score for class c on input X[i].

If y is not None, instead return a tuple of:

- loss: Loss (data loss and regularization loss) for this batch of training samples.
- grads: Dictionary mapping parameter names to gradients of those parameters
with respect to the loss function; has the same keys as self.params.

"""

```
# Unpack variables from the params dictionary  
W1, b1 = self.params['W1'], self.params['b1']  
W2, b2 = self.params['W2'], self.params['b2']
```

```

N, D = X.shape

# Compute the forward pass
scores = None

# =====
#
# YOUR CODE HERE:
#   Calculate the output scores of the neural network. The
result      #   should be (N, C). As stated in the description for this
class,      #   there should not be a ReLU layer after the second FC
layer.      #   The output of the second FC layer is the output
scores. Do not use a for loop in your implementation.
# =====
#

f = lambda x : x * (x > 0)
a = X.dot(W1.T) + b1
h1 = f(a)
scores = h1.dot(W2.T) + b2

# =====
#
# END YOUR CODE HERE
# =====
#

# If the targets are not given then jump out, we're done
if y is None:
    return scores

# Compute the loss
loss = None

# =====
#
# YOUR CODE HERE:
#   Calculate the loss of the neural network. This includes
the         #   softmax loss and the L2 regularization for W1 and W2.
Store the   #   total loss in the variable loss. Multiply the
regularization #   loss by 0.5 (in addition to the factor reg).
#
#

```

```

===== #

num_train = X.shape[0]
selector = np.arange(num_train), y

temp_scores = scores
temp_scores -= np.max(temp_scores)
loss = np.sum(np.log(np.sum(np.exp(temp_scores).T, axis=0)) -
temp_scores[selector]) / num_train
loss += 0.5 * reg * (np.sum(W1**2) + np.sum(W2**2))

# scores is num_examples by num_classes

# =====
# # END YOUR CODE HERE
# =====
#
grads = {}

# =====
# # YOUR CODE HERE:
#       #       Implement the backward pass. Compute the derivatives
of the      #       weights and the biases. Store the results in the
grads       #       dictionary. e.g., grads['W1'] should store the
gradient for #       W1, and be of the same size as W1.
#
===== #

prob = np.exp(scores)
row_sums = np.sum(prob, axis=1)
prob = prob / row_sums[:, np.newaxis]
prob[range(prob.shape[0]), y] -= 1
prob /= X.shape[0]

dldh2 = prob.sum(axis=0)

grads['b2'] = dldh2
grads['W2'] = h1.T.dot(prob).T + reg * W2

dldh1 = prob.dot(W2)
dllda = dldh1 * (a > 0)

grads['b1'] = dllda.sum(axis=0)
grads['W1'] = X.T.dot(dllda).T + reg * W1

```

```

# =====
#
# END YOUR CODE HERE
# =====
#

return loss, grads

def train(self, X, y, X_val, y_val,
          learning_rate=1e-3, learning_rate_decay=0.95,
          reg=1e-5, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this neural network using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c
    means that X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the
    learning rate after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]
    iterations_per_epoch = max(num_train / batch_size, 1)

    # Use SGD to optimize the parameters in self.model
    loss_history = []
    train_acc_history = []
    val_acc_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        #
        ===== #
        # YOUR CODE HERE:
        #     Create a minibatch by sampling batch_size samples
        randomly.
        #
        ===== #

```

```

indices = np.random.choice(num_train, batch_size)
X_batch, y_batch = X[indices], y[indices]

#
===== #
# END YOUR CODE HERE
#
===== #

# Compute loss and gradients using the current minibatch
loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
loss_history.append(loss)

#
===== #
# YOUR CODE HERE:
#     Perform a gradient descent step using the minibatch
to update
#     all parameters (i.e., W1, W2, b1, and b2).
#
===== #

loss, grads = self.loss(X_batch, y_batch, reg=reg)
loss_history.append(loss)

self.params['W1'] -= learning_rate * grads['W1']
self.params['b1'] -= learning_rate * grads['b1']
self.params['W2'] -= learning_rate * grads['W2']
self.params['b2'] -= learning_rate * grads['b2']

#
===== #
# END YOUR CODE HERE
#
===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))

# Every epoch, check train and val accuracy and decay learning
rate.
if it % iterations_per_epoch == 0:
    # Check accuracy
    train_acc = (self.predict(X_batch) == y_batch).mean()
    val_acc = (self.predict(X_val) == y_val).mean()
    train_acc_history.append(train_acc)
    val_acc_history.append(val_acc)

    # Decay learning rate

```

```

        learning_rate *= learning_rate_decay

    return {
        'loss_history': loss_history,
        'train_acc_history': train_acc_history,
        'val_acc_history': val_acc_history,
    }

    def predict(self, X):
        """
        Use the trained weights of this two-layer network to predict
        labels for
        data points. For each data point we predict scores for each of the
        C
        classes, and assign each data point to the class with the highest
        score.

        Inputs:
        - X: A numpy array of shape (N, D) giving N D-dimensional data
        points to
        classify.

        Returns:
        - y_pred: A numpy array of shape (N,) giving predicted labels for
        each of
        the elements of X. For all i, y_pred[i] = c means that X[i] is
        predicted
        to have class c, where 0 <= c < C.
        """
        y_pred = None

        # =====
        #
        # YOUR CODE HERE:
        #   Predict the class given the input data.
        # =====
        #
        f = lambda x : x * (x > 0)
        h1 = f(X.dot(self.params['W1']).T) + self.params['b1']
        scores = h1.dot(self.params['W2']).T + self.params['b2']
        y_pred = scores.argmax(axis=1)

        # =====
        #
        # END YOUR CODE HERE
        # =====
        #

    return y_pred

```

