

```

import numpy as np
import pdb

from .layers import *
from .layer_utils import *

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu). It has been modified in various areas for use
in the
ECE 239AS class at UCLA. This includes the descriptions of what code
to
implement as well as some slight potential changes in variable names
to be
consistent with class nomenclature. We thank Justin Johnson & Serena
Yeung for
permission to use this code. To see the original version, please
visit
cs231n.stanford.edu.
"""

class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network with ReLU nonlinearity
    and
    softmax loss that uses a modular layer design. We assume an input
    dimension
    of D, a hidden dimension of H, and perform classification over C
    classes.

    The architecture should be affine - relu - affine - softmax.

    Note that this class does not implement gradient descent; instead,
    it
    will interact with a separate Solver object that is responsible
    for running
    optimization.

    The learnable parameters of the model are stored in the dictionary
    self.params that maps parameter names to numpy arrays.
    """

    def __init__(self, input_dim=3*32*32, hidden_dims=100,
num_classes=10,
        dropout=0, weight_scale=1e-3, reg=0.0):
        """
        Initialize a new network.

        Inputs:
        - input_dim: An integer giving the size of the input

```

```

        - hidden_dims: An integer giving the size of the hidden layer
        - num_classes: An integer giving the number of classes to
classify
        - dropout: Scalar between 0 and 1 giving dropout strength.
        - weight_scale: Scalar giving the standard deviation for
random
        initialization of the weights.
        - reg: Scalar giving L2 regularization strength.
        """
        self.params = {}
        self.reg = reg

        #
===== #
        # YOUR CODE HERE:
        # Initialize W1, W2, b1, and b2. Store these as
self.params['W1'],
        # self.params['W2'], self.params['b1'] and
self.params['b2']. The
        # biases are initialized to zero and the weights are
initialized
        # so that each parameter has mean 0 and standard deviation
weight_scale.
        # The dimensions of W1 should be (input_dim, hidden_dim) and
the
        # dimensions of W2 should be (hidden_dims, num_classes)
        #
===== #

        self.params['W1'] = np.random.normal(scale=weight_scale,
size=(input_dim, hidden_dims))
        self.params['b1'] = np.zeros(hidden_dims)
        self.params['W2'] = np.random.normal(scale=weight_scale,
size=(hidden_dims, num_classes))
        self.params['b2'] = np.zeros(num_classes)

        #
===== #
        # END YOUR CODE HERE
        #
===== #

def loss(self, X, y=None):
    """
    Compute loss and gradient for a minibatch of data.

    Inputs:
    - X: Array of input data of shape (N, d_1, ..., d_k)
    - y: Array of labels, of shape (N,). y[i] gives the label for
X[i].

```

```

        Returns:
        If y is None, then run a test-time forward pass of the model
and return:
    - scores: Array of shape (N, C) giving classification scores,
where
    scores[i, c] is the classification score for X[i] and class
c.

    If y is not None, then run a training-time forward and
backward pass and
    return a tuple of:
    - loss: Scalar value giving the loss
    - grads: Dictionary with the same keys as self.params, mapping
parameter
    names to gradients of the loss with respect to those
parameters.
    """
    scores = None

    #
===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the two-layer neural
network. Store
    # the class scores as the variable 'scores'. Be sure to use
the layers
    # you prior implemented.
    #
===== #

    hidden, cache1 = affine_relu_forward(X, self.params["W1"],
self.params["b1"])
    scores, cache2 = affine_forward(hidden, self.params["W2"],
self.params["b2"])

    #
===== #
    # END YOUR CODE HERE
    #
===== #

    # If y is None then we are in test mode so just return scores
if y is None:
    return scores

    loss, grads = 0, {}
    #
===== #
    # YOUR CODE HERE:

```

```

        # Implement the backward pass of the two-layer neural net.
Store   # the loss as the variable 'loss' and store the gradients in
the     # 'grads' dictionary. For the grads dictionary, grads['W1']
holds   # the gradient for W1, grads['b1'] holds the gradient for
        # b1, etc.
        # i.e., grads[k] holds the gradient for self.params[k].
        #
        # Add L2 regularization, where there is an added cost
0.5*self.reg*W^2
to       # for each W. Be sure to include the 0.5 multiplying factor
        # match our implementation.
        #
        # And be sure to use the layers you prior implemented.
        #
===== #

        l, g = softmax_loss(scores, y)
        reg_loss = 0.5 * self.reg * (np.sum(self.params["W1"]**2) +
np.sum(self.params["W2"]**2))

        dldx2, grads["W2"], grads["b2"] = affine_backward(g, cache2)
        dldx1, grads["W1"], grads["b1"] = affine_relu_backward(dldx2,
cache1)

        loss = l + reg_loss
        grads["W1"] += self.reg * self.params["W1"]
        grads["W2"] += self.reg * self.params["W2"]

        #
===== #
        # END YOUR CODE HERE
        #
===== #

    return loss, grads

```

```

class FullyConnectedNet(object):
    """

```

```

    A fully-connected neural network with an arbitrary number of
hidden layers,
    ReLU nonlinearities, and a softmax loss function. This will also
implement
    dropout and batch normalization as options. For a network with L
layers,
    the architecture will be

```

`{affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax`

where batch normalization and dropout are optional, and the {...} block is repeated $L - 1$ times.

Similar to the TwoLayerNet above, learnable parameters are stored in the `self.params` dictionary and will be learned using the Solver class.

```
def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
              dropout=1, use_batchnorm=False, reg=0.0,
              weight_scale=1e-2, dtype=np.float32, seed=None):
    """
    Initialize a new FullyConnectedNet.

    Inputs:
    - hidden_dims: A list of integers giving the size of each
hidden layer.
    - input_dim: An integer giving the size of the input.
    - num_classes: An integer giving the number of classes to
classify.
    - dropout: Scalar between 0 and 1 giving dropout strength. If
dropout=1 then
        the network should not use dropout at all.
    - use_batchnorm: Whether or not the network should use batch
normalization.
    - reg: Scalar giving L2 regularization strength.
    - weight_scale: Scalar giving the standard deviation for
random
        initialization of the weights.
    - dtype: A numpy datatype object; all computations will be
performed using
        this datatype. float32 is faster but less accurate, so you
should use
        float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout
layers. This
        will make the dropout layers deterministic so we can
gradient check the
        model.
    """
    self.use_batchnorm = use_batchnorm
    self.use_dropout = dropout < 1
    self.reg = reg
    self.num_layers = 1 + len(hidden_dims)
    self.dtype = dtype
```

```

self.params = {}

#
===== #
# YOUR CODE HERE:
# Initialize all parameters of the network in the
self.params dictionary.
# The weights and biases of layer 1 are W1 and b1; and in
general the
# weights and biases of layer i are Wi and bi. The
# biases are initialized to zero and the weights are
initialized
# so that each parameter has mean 0 and standard deviation
weight_scale.
#
# BATCHNORM: Initialize the gammas of each layer to 1 and
the beta
# parameters to zero. The gamma and beta parameters for
layer 1 should
# be self.params['gamma1'] and self.params['beta1']. For
layer 2, they
# should be gamma2 and beta2, etc. Only use batchnorm if
self.use_batchnorm
# is true and DO NOT do batch normalize the output scores.
#
===== #

n = len(hidden_dims)
self.params["W1"] = np.random.normal(scale=weight_scale,
size=(input_dim, hidden_dims[0]))
self.params["b1"] = np.zeros(hidden_dims[0])

for i in range(1, n):
    self.params["W" + str(i+1)] =
np.random.normal(scale=weight_scale, size=(hidden_dims[i-1],
hidden_dims[i]))
    self.params["b" + str(i+1)] = np.zeros(hidden_dims[i])

self.params["W" + str(n+1)] =
np.random.normal(scale=weight_scale, size=(hidden_dims[n-1],
num_classes))
self.params["b" + str(n+1)] = np.zeros(num_classes)

if self.use_batchnorm:
    for i in range(self.num_layers - 1):
        self.params["gamma" + str(i+1)] =
np.ones(hidden_dims[i])
        self.params["beta" + str(i+1)] =
np.zeros(hidden_dims[i])

```

```

#
===== #
# END YOUR CODE HERE
#
===== #

# When using dropout we need to pass a dropout_param
dictionary to each
# dropout layer so that the layer knows the dropout
probability and the mode
# (train / test). You can pass the same dropout_param to each
dropout layer.
self.dropout_param = {}
if self.use_dropout:
    self.dropout_param = {'mode': 'train', 'p': dropout}
if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running
means and
# variances, so we need to pass a special bn_param object to
each batch
# normalization layer. You should pass self.bn_params[0] to
the forward pass
# of the first batch normalization layer, self.bn_params[1] to
the forward
# pass of the second batch normalization layer, etc.
self.bn_params = []
if self.use_batchnorm:
    self.bn_params = [{'mode': 'train'} for i in
np.arange(self.num_layers - 1)]

# Cast all parameters to the correct datatype
for k, v in self.params.items():
    self.params[k] = v.astype(dtype)

def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param
since they
# behave differently during training and testing.
if self.dropout_param is not None:

```

```

        self.dropout_param['mode'] = mode
    if self.use_batchnorm:
        for bn_param in self.bn_params:
            bn_param[mode] = mode

    scores = None

    #
    ===== #
    # YOUR CODE HERE:
    # Implement the forward pass of the FC net and store the
output    # scores as the variable "scores".
    #
    # BATCHNORM: If self.use_batchnorm is true, insert a
batchnorm    # layer
    # between the affine_forward and relu_forward layers. You
may    #
    # also write an affine_batchnorm_relu() function in
layer_utils.py.
    #
    # DROPOUT: If dropout is non-zero, insert a dropout layer
after    #
    # every ReLU layer.
    #
    ===== #

    caches = []
    dp_caches = []
    fwd_in = X

    for i in range(self.num_layers - 1):
        if self.use_batchnorm:
            fwd_in, cache = affine_batchnorm_relu_forward(fwd_in,
+ str(i+1)],                                     self.params["W"
+ str(i+1)],                                     self.params["b"
self.params["gamma" + str(i+1)],
self.params["beta" + str(i+1)],
self.bn_params[i])
            caches.append(cache)
        else:
            fwd_in, cache = affine_relu_forward(fwd_in,
self.params["W" + str(i+1)], self.params["b" + str(i+1)])
            caches.append(cache)
        if self.use_dropout:

```



```

        fwd_in, dp_cache = dropout_forward(fwd_in,
self.dropout_param)
        dp_caches.append(dp_cache)

        scores, cache = affine_forward(fwd_in, self.params["W" +
str(self.num_layers)], self.params["b" + str(self.num_layers)])
        caches.append(cache)

#
===== #
# END YOUR CODE HERE
#
===== #

# If test mode return early
if mode == 'test':
    return scores

loss, grads = 0.0, {}
#
===== #
# YOUR CODE HERE:
#   Implement the backwards pass of the FC net and store the
gradients
#   in the grads dict, so that grads[k] is the gradient of
self.params[k]
#   Be sure your L2 regularization includes a 0.5 factor.
#
#   BATCHNORM: Incorporate the backward pass of the batchnorm.
#
#   DROPOUT: Incorporate the backward pass of dropout.
#
===== #

l, g = softmax_loss(scores, y)
reg_loss = 0.5 * self.reg * (sum([np.sum(self.params["W" +
str(i)]**2) for i in range(1, self.num_layers + 1)]))

upstream_deriv = g

upstream_deriv, grads["W" + str(self.num_layers)], grads["b" +
str(self.num_layers)] = affine_backward(upstream_deriv,
caches[self.num_layers - 1])

for i in range(self.num_layers - 2, -1, -1):
    if self.use_dropout:
        upstream_deriv = dropout_backward(upstream_deriv,
dp_caches[i])
    if self.use_batchnorm:
        upstream_deriv, grads["W" + str(i+1)], grads["b" +

```

```

str(i+1)], grads["gamma" + str(i+1)], grads["beta" + str(i+1)] =
affine_batchnorm_relu_backward(upstream_deriv, caches[i])
    else:
        upstream_deriv, grads["W" + str(i+1)], grads["b" +
str(i+1)] = affine_relu_backward(upstream_deriv, caches[i])

    loss = l + reg_loss
    for i in range(self.num_layers):
        grads["W" + str(i+1)] += self.reg * self.params["W" +
str(i+1)]

#
===== #
# END YOUR CODE HERE
#
===== #

    return loss, grads

```