

```

import numpy as np
from nndl.layers import *
import pdb

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu). It has been modified in various areas for use
in the
ECE 239AS class at UCLA. This includes the descriptions of what code
to
implement as well as some slight potential changes in variable names
to be
consistent with class nomenclature. We thank Justin Johnson & Serena
Yeung for
permission to use this code. To see the original version, please
visit
cs231n.stanford.edu.
"""

def conv_forward_naive(x, w, b, conv_param):
    """
    A naive implementation of the forward pass for a convolutional
    layer.

    The input consists of N data points, each with C channels, height H
    and width
    W. We convolve each input with F different filters, where each
    filter spans
    all C channels and has height HH and width HH.

    Input:
    - x: Input data of shape (N, C, H, W)
    - w: Filter weights of shape (F, C, HH, WW)
    - b: Biases, of shape (F,)
    - conv_param: A dictionary with the following keys:
        - 'stride': The number of pixels between adjacent receptive fields
        in the
        horizontal and vertical directions.
        - 'pad': The number of pixels that will be used to zero-pad the
        input.

    Returns a tuple of:
    - out: Output data, of shape (N, F, H', W') where H' and W' are
    given by
         $H' = 1 + (H + 2 * \text{pad} - \text{HH}) / \text{stride}$ 
         $W' = 1 + (W + 2 * \text{pad} - \text{WW}) / \text{stride}$ 
    - cache: (x, w, b, conv_param)
    """
    out = None
    pad = conv_param['pad']

```

```

stride = conv_param['stride']

# ===== #
# YOUR CODE HERE:
# Implement the forward pass of a convolutional neural network.
# Store the output as 'out'.
# Hint: to pad the array, you can use the function np.pad.
# ===== #

n, c, xh, xw = x.shape
wf, wc, wh, ww = w.shape
hout, wout = 1 + (xh + 2*pad - wh)//stride, 1 + (xw + 2*pad - ww)//
stride

out = np.zeros((len(x), wf, hout, wout))

# pad the input
x_padded = np.pad(x, ((0, 0), (0, 0), (pad, pad), (pad, pad)),
'constant', constant_values=0)

# loop over all inputs
for i in range(n):
    # loop over filters
    for j in range(wf):
        # loop over each cell in output
        for hi in range(hout):
            for wi in range(wout):
                out[i, j, hi, wi] = np.sum(x_padded[i, :, hi*stride:
hi*stride + wh, wi*stride: wi*stride + ww] * w[j]) + b[j]

# ===== #
# END YOUR CODE HERE
# ===== #

cache = (x, w, b, conv_param)
return out, cache

```

```

def conv_backward_naive(dout, cache):
    """

```

A naive implementation of the backward pass for a convolutional layer.

Inputs:

- dout: Upstream derivatives.
- cache: A tuple of (x, w, b, conv\_param) as in conv\_forward\_naive

Returns a tuple of:

- dx: Gradient with respect to x
- dw: Gradient with respect to w

```

- db: Gradient with respect to b
"""
dx, dw, db = None, None, None

N, F, out_height, out_width = dout.shape
x, w, b, conv_param = cache

stride, pad = [conv_param['stride'], conv_param['pad']]
xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)),
mode='constant')
num_filts, _, f_height, f_width = w.shape

# ===== #
# YOUR CODE HERE:
# Implement the backward pass of a convolutional neural network.
# Calculate the gradients: dx, dw, and db.
# ===== #

xn, xc, xh, xw = xpad.shape
d_n, d_f, d_h, d_w = dout.shape
wf, wc, wh, ww = w.shape

dw = np.zeros(w.shape)
dx = np.zeros(xpad.shape)
db = np.zeros(b.shape)

for i in range(d_n):
    for j in range(d_f):
        for k in range(d_h):
            for l in range(d_w):
                dw[j] += xpad[i, :, k*stride: k*stride + wh,
l*stride: l*stride + ww] * dout[i, j, k, l]
                db[j] += dout[i, j, k, l]
                dx[i, :, k*stride: k*stride + wh, l*stride: l*stride
+ ww] += dout[i, j, k, l] * w[j]

# the pad is used to help calculate the gradient, so we now remove
the pad from the gradient
xpad_h, xpad_w = xpad.shape[2], xpad.shape[3]
dx = dx[:, :, pad:xpad_h - pad, pad:xpad_w - pad]

# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dw, db

def max_pool_forward_naive(x, pool_param):
    """

```

A naive implementation of the forward pass for a max pooling layer.

Inputs:

- x: Input data, of shape (N, C, H, W)
- pool\_param: dictionary with the following keys:
  - 'pool\_height': The height of each pooling region
  - 'pool\_width': The width of each pooling region
  - 'stride': The distance between adjacent pooling regions

Returns a tuple of:

- out: Output data
- cache: (x, pool\_param)

"""

out = None

```
# ===== #  
# YOUR CODE HERE:  
#   Implement the max pooling forward pass.  
# ===== #
```

```
p_h, p_w, stride = pool_param['pool_height'],  
pool_param['pool_width'], pool_param['stride']  
n, c, h, w = x.shape  
h_new, w_new = (h - p_h)//stride + 1, (w - p_w)//stride + 1  
  
out = np.zeros((n, c, h_new, w_new))  
  
for i in range(n):  
    for j in range(c):  
        for k in range(h_new):  
            for l in range(w_new):  
                out[i, j, k, l] = np.max(x[i, j, k*stride: k*stride  
+ p_h, l*stride: l*stride + p_w])
```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
cache = (x, pool_param)  
return out, cache
```

```
def max_pool_backward_naive(dout, cache):  
    """
```

A naive implementation of the backward pass for a max pooling layer.

Inputs:

- dout: Upstream derivatives
- cache: A tuple of (x, pool\_param) as in the forward pass.

Returns:

- dx: Gradient with respect to x

```

"""
dx = None
x, pool_param = cache
pool_height, pool_width, stride = pool_param['pool_height'],
pool_param['pool_width'], pool_param['stride']

# ===== #
# YOUR CODE HERE:
#   Implement the max pooling backward pass.
# ===== #

dout_n, dout_c, dout_w, dout_h = dout.shape

dx = np.zeros(x.shape)

for i in range(dout_n):
    for j in range(dout_c):
        for k in range(dout_w):
            for l in range(dout_h):
                slice = x[i, j, k*stride: k*stride + pool_height,
l*stride: l*stride + pool_width]
                mask = slice == np.max(slice)
                dx[i, j, k*stride: k*stride + pool_height, l*stride:
l*stride + pool_width] += mask * dout[i, j, k, l]

# ===== #
# END YOUR CODE HERE
# ===== #

return dx

```

```

def spatial_batchnorm_forward(x, gamma, beta, bn_param):
    """

```

Computes the forward pass for spatial batch normalization.

Inputs:

- x: Input data of shape (N, C, H, W)
- gamma: Scale parameter, of shape (C,)
- beta: Shift parameter, of shape (C,)
- bn\_param: Dictionary with the following keys:
  - mode: 'train' or 'test'; required
  - eps: Constant for numeric stability
  - momentum: Constant for running mean / variance. momentum=0 means

that

old information is discarded completely at every time step,

while

momentum=1 means that new information is never incorporated. The default of momentum=0.9 should work well in most situations.

- running\_mean: Array of shape (D,) giving running mean of features

- running\_var Array of shape (D,) giving running variance of features

Returns a tuple of:

- out: Output data, of shape (N, C, H, W)
- cache: Values needed for the backward pass

"""

out, cache = None, None

# ===== #

# YOUR CODE HERE:

# Implement the spatial batchnorm forward pass.

#

# You may find it useful to use the batchnorm forward pass you

# implemented in HW #4.

# ===== #

n, c, h, w = x.shape

x = x.transpose(0, 2, 3, 1).reshape((n\*h\*w, c))

out, cache = batchnorm\_forward(x, gamma, beta, bn\_param)

out = out.reshape(n, h, w, c).transpose(0, 3, 1, 2)

# ===== #

# END YOUR CODE HERE

# ===== #

return out, cache

def spatial\_batchnorm\_backward(dout, cache):

"""

Computes the backward pass for spatial batch normalization.

Inputs:

- dout: Upstream derivatives, of shape (N, C, H, W)
- cache: Values from the forward pass

Returns a tuple of:

- dx: Gradient with respect to inputs, of shape (N, C, H, W)
- dgamma: Gradient with respect to scale parameter, of shape (C,)
- dbeta: Gradient with respect to shift parameter, of shape (C,)

"""

dx, dgamma, dbeta = None, None, None

# ===== #

# YOUR CODE HERE:

# Implement the spatial batchnorm backward pass.

#

# You may find it useful to use the batchnorm backward pass you

# implemented in HW #4.

```
# ===== #
n, c, h, w = dout.shape
dout = dout.transpose(0, 2, 3, 1).reshape((n*h*w, c))
dx, dgamma, dbeta = batchnorm_backward(dout, cache)
dx = dx.reshape(n, h, w, c).transpose(0, 3, 1, 2)

# ===== #
# END YOUR CODE HERE
# ===== #

return dx, dgamma, dbeta
```